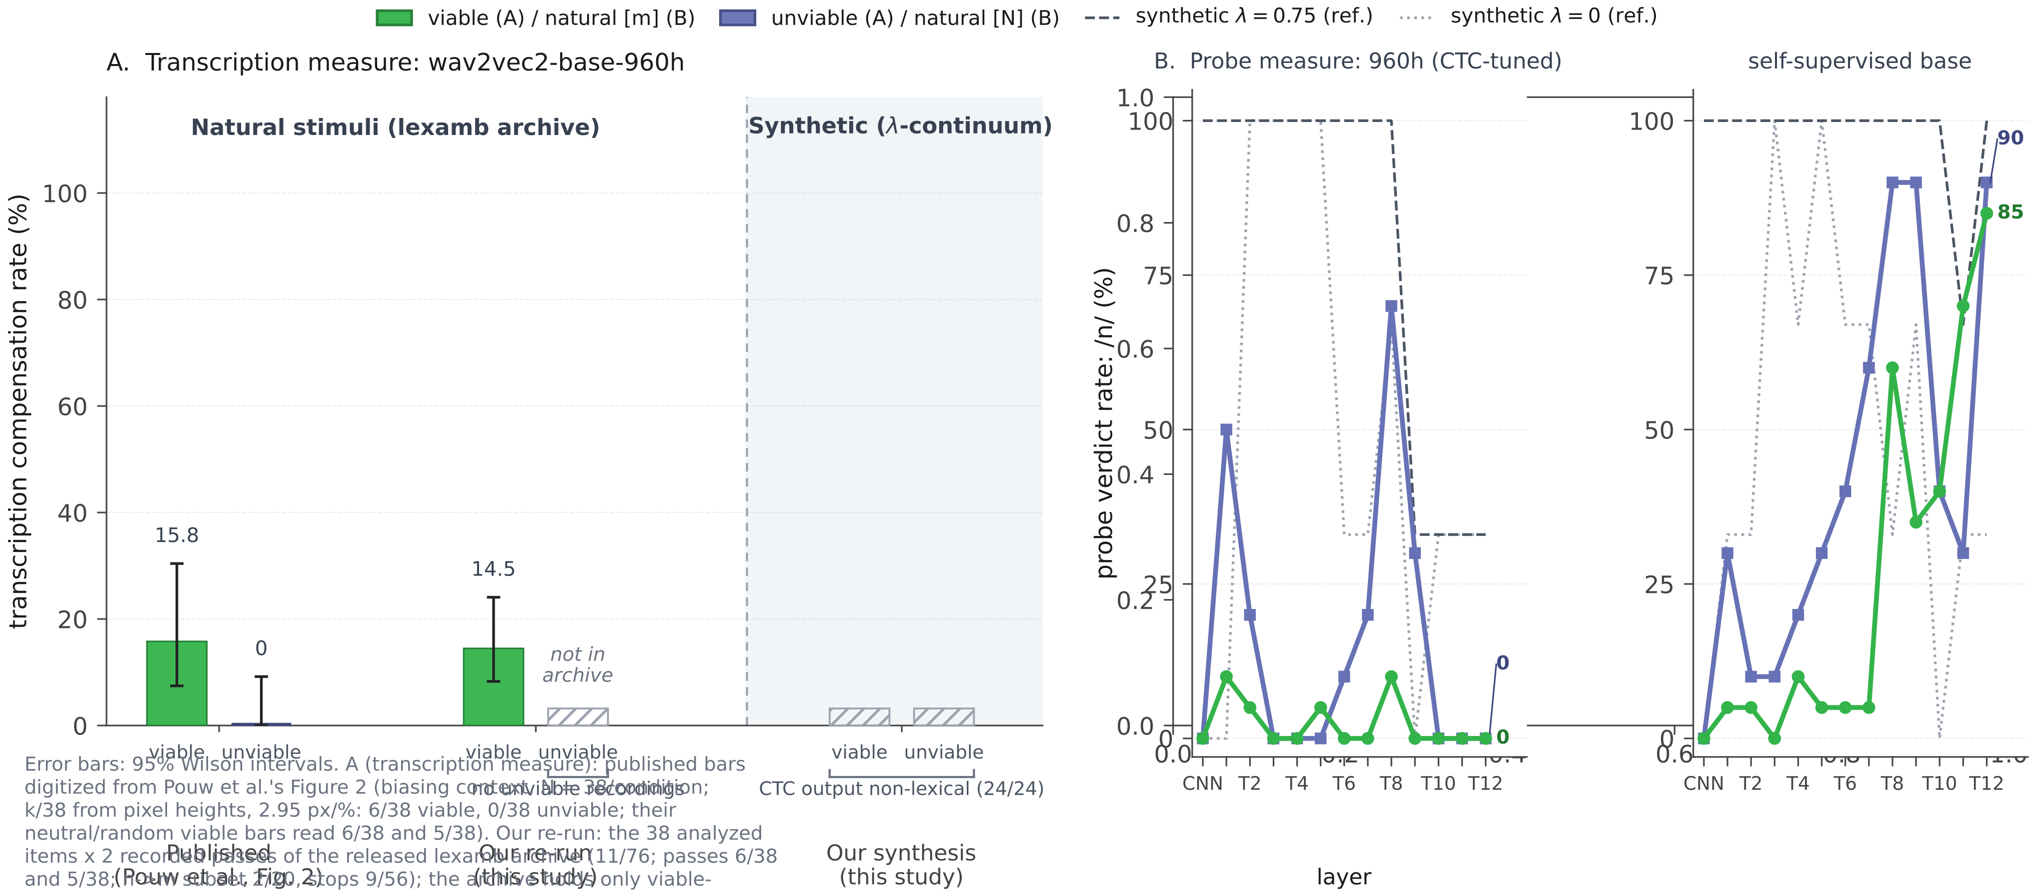

Error bars: 95% Wilson intervals. A (transcription measure): published bars digitized from Pouw et al.'s Figure 2 (biasing context, N = 38/condition; k/38 from pixel heights, 2.95 px/%: 6/38 viable, 0/38 unviable; their neutral/random viable bars read 6/38 and 5/38). Our re-run: the 38 analyzed items x 2 recorded passes of the released lexamb archive (11/76; passes 6/38 and 5/38; n->m subset 2/20, stops 9/56); the archive holds only viable-context recordings, and the same ASR head returns non-lexical transcriptions on our VocalTractLab continuum (24/24) - hence the probe measure on the synthetic side. B (probe measure): rate of verdict /n/ per layer; the same 13 layer probes as every experiment here, trained on the 18 canonical training sentences. Natural [m]: 10 n->m items x 2 passes (n = 20); natural [N]: the 5 items excluded by Pouw et al.'s transcription measure (n = 10, exploratory). Synthetic references pool the 3 licensed frames per layer.